




EXECUTIVE SUMMARY

Being findable and being citable are now different problems

Ranking #1 on a search results page used to be enough, a human would scroll to your link. An answer engine skips the page entirely: it retrieves a handful of passages, reasons over them, and writes one answer. Your brand either enters that working set and gets named, or it doesn't exist for that query.

Sections 1–3 define AI visibility as a measurable quantity. Section 4 models it as a logistic function of six evidence factors and *proves* the function is strictly increasing in each one (Theorem 1), with a concavity corollary explaining why a spread portfolio of improvements beats maxing out a single lever. Section 7 turns this from theory into a measurement protocol, including a control-group design for isolating causal effect from background noise.

Core objective

$$\max_B \mathbb{E}_{q \sim Q} [P_B(q)]$$

Maximize the expected probability that the AI includes brand B across the query universe Q a customer actually asks.

1. DEFINING AI VISIBILITY

A brand either appears in the answer, or it doesn't

Fix a brand B and a finite set of commercially relevant questions, the *query universe*:

$$Q = \{q_1, q_2, \dots, q_n\}$$

e.g. "best project management software for agencies," "alternatives to Asana," "how should an agency manage client projects?", the informational, commercial, comparative and problem-oriented questions a real buyer asks before purchasing. Define the appearance indicator:

$$V(B, q) = \begin{cases} 1 & \text{if } B \text{ appears in the AI answer to } q \\ 0 & \text{otherwise} \end{cases}$$

The simplest visibility statistic is the sample mean of this indicator over the query universe:

DEFINITION 1: AI VISIBILITY RATE

$$AIV(B) = \frac{1}{n} \sum_{i=1}^n V(B, q_i)$$

If B appears in 37 of 100 representative queries, $AIV(B) = 0.37$. This is already a measurable KPI, Section 6 treats it as a statistical estimator with a confidence interval, not just a percentage.

2. THREE DIMENSIONS OF APPEARING

Mentioned, recommended, and cited are not the same event

Refine the single indicator V into three binary sub-indicators of increasing commercial value:

$$M(B, q), R_{\text{ec}}(B, q), C(B, q) \in \{0, 1\}$$

these are mentioned, recommended, and cited (backed by a named source), respectively, and combine them into a weighted score:

$$S(B, q) = w_m M(B, q) + w_r R_{\text{ec}}(B, q) + w_c C(B, q), \quad w_m + w_r + w_c = 1$$

With, say, $w_m = 0.2$, $w_r = 0.4$, $w_c = 0.4$: a bare mention scores 0.2; a recommendation scores 0.4; a recommendation backed by a citation scores 0.8. This is a considerably more useful signal than a binary "we showed up", it separates being named from being endorsed.

3. THE OBJECTIVE AS AN EXPECTATION

Treat appearance as a probability, not a binary fact

A given brand doesn't deterministically appear or not for a given query, sampling temperature, retrieval variance, and index drift make it a random event. Define:

$$P_B(q) := \Pr(B \text{ appears in the answer} \mid q)$$

and the optimization target is the expectation of this probability over the query universe a real buyer asks:

THE OBJECTIVE

$$\max_B \mathbb{E}_{q \sim Q} [P_B(q)]$$

In plain language: maximize the probability that the AI includes your brand, averaged over the commercially relevant questions people actually ask. Section 4 asks what $P_B(q)$ is a function of.

4. THE EVIDENCE MODEL

Theorem 1: citation probability is monotonic in every lever

No AI search provider publishes its ranking function, and this is not a claim to have reverse engineered one. It is a stylized, testable model, a logistic function of six factors known to matter to retrieval-and-generation systems generally, built so its qualitative predictions (monotonicity, diminishing returns) hold for a wide family of coefficient values, not one guessed set of weights.

DEFINITION 2: THE SIX EVIDENCE FACTORS

Let $R, A, E, T, S, F \in [0, 1]$ denote, respectively, a brand's relevance to the query, authority, independent evidence, topical coverage, semantic/entity clarity, and content freshness. Define the linear predictor and the logistic citation probability:

$$z_B(q) = \beta_0 + \beta_1 R + \beta_2 A + \beta_3 E + \beta_4 T + \beta_5 S + \beta_6 F$$

$$P_B(q) = \sigma(z_B(q)), \quad \sigma(x) = \frac{1}{1 + e^{-x}}$$

THEOREM 1: MONOTONICITY

If $\beta_k > 0$ for every $k = 1, \dots, 6$, i.e. each factor genuinely helps and none hurts, then for every factor $x_k \in \{R, A, E, T, S, F\}$:

$$\frac{\partial P_B}{\partial x_k} = \beta_k \sigma(z_B) (1 - \sigma(z_B)) > 0$$

Proof. The logistic function satisfies $\sigma'(z) = \sigma(z)(1 - \sigma(z))$, which is strictly positive for every finite z since $\sigma(z) \in (0, 1)$. By the chain rule, $\partial P_B / \partial x_k = \beta_k \cdot \sigma'(z_B)$. Since $\beta_k > 0$ by hypothesis and $\sigma'(z_B) > 0$ always, the product is strictly positive. ■

Every lever that genuinely carries positive weight strictly raises citation probability, none of the six factors can be safely ignored, and none can hurt you by being improved.

COROLLARY 1: DIMINISHING RETURNS, AND WHY TO DIVERSIFY

Differentiating again, $\sigma''(z) = \sigma'(z)(1 - 2\sigma(z))$, which changes sign at $\sigma(z) = \frac{1}{2}$ (i.e. $z = 0$). For $z > 0$, $\sigma''(z) < 0$: the marginal gain $\sigma'(z)$ is *falling* as z grows. Since $\partial^2 P_B / \partial x_k^2 = \beta_k^2 \sigma''(z_B)$ has the same sign as $\sigma''(z_B)$, P_B is concave in each x_k once $z_B > 0$.

$$\frac{\partial^2 P_B}{\partial x_k^2} = \beta_k^2 \sigma'(z_B)(1 - 2\sigma(z_B)) < 0 \quad \text{once } z_B > 0$$

A budget-constrained optimum equalizes marginal effect $\beta_k \sigma'(z_B)$ across the six levers rather than exhausting it on whichever is already strongest. In practice: once one lever (say authority) has already pushed z_B well past zero, another unit of authority buys less citation probability than the same unit spent raising a neglected lever (say freshness), the formal reason a diversified GEO program outperforms a single-tactic one.

5. EVIDENCE ACCUMULATION

Why one page you wrote is worth less than five independent ones

The authority and evidence factors above (A, E) are themselves built from many independent sources. Model each source i as contributing evidence e_i , weighted by its own credibility w_i , into a pooled evidence score, which crosses a threshold θ before the system is confident enough to recommend:

$$E_B = \sum_{i=1}^k w_i e_i, \quad \Pr(\text{recommend} \mid E_B) = \sigma(E_B - \theta)$$

A weak brand might sit at $E_B = 1.2$, an established one at $E_B = 7.8$, a category leader at $E_B = 15.4$, as independent, credible evidence accumulates, $\sigma(E_B - \theta)$ rises toward 1. Treated as a graph, brand, product, category, publication, and customer nodes, connected by "mentions" and "reviews" edges, the practical goal is a *dense, consistent subgraph* around the brand:

MODELING HEURISTIC, NOT A DERIVED IDENTITY

Brand strength $\propto$ Authority $\times$ Relevance $\times$ Evidence density $\times$ Consistency

Unlike Theorem 1, this proportionality is a conceptual heuristic for why distributed third-party evidence compounds rather than adds, it is not proved from first principles here. It explains, informally, why 500 brand-authored blog posts move the needle far less than 20 independent ones: they add to only one node's out-degree, not to the density of the surrounding subgraph.

6. MEASURING IT

From estimator to causal proof

6.1: POINT ESTIMATE AND CONFIDENCE INTERVAL

Sampling n queries and counting x appearances gives the estimator $\hat{p} = x/n$ for $AIV(B)$, with the usual normal-approximation interval:

$$\hat{p} = \frac{x}{n}, \quad \hat{p} \pm z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

6.2: BEFORE / AFTER LIFT

With $n = 1,000$ prompts, $x = 120$ mentions before optimization gives $\hat{p}_{\text{before}} = 12\%$; $x = 230$ after gives $\hat{p}_{\text{after}} = 23\%$:

$$\Delta = \hat{p}_{\text{after}} - \hat{p}_{\text{before}} = 11 \text{ pp}, \quad \frac{\Delta}{\hat{p}_{\text{before}}} = \frac{11}{12} \approx 91.7\% \text{ relative lift}$$

Measurable, but not yet *causal*. AI visibility can drift for everyone as models and indices change; a single before/after comparison can't separate your intervention from that drift.

6.3: DIFFERENCE-IN-DIFFERENCES, ISOLATING THE CAUSAL EFFECT

Optimize 50 topics; hold out 50 comparable topics as an unoptimized control:

	Before	After
Optimized topics	12%	28%
Control topics	11%	14%

$$\widehat{\text{DiD}} = (\hat{p}_{T,\text{after}} - \hat{p}_{T,\text{before}}) - (\hat{p}_{C,\text{after}} - \hat{p}_{C,\text{before}}) = (28 - 12) - (14 - 11) = \boxed{13 \text{ pp}}$$

The control group's 3-point drift is background noise (new model version, seasonal demand, whatever). Subtracting it leaves 13 points attributable to the intervention, far stronger evidence than "our AI visibility went up."

6.4: A COMPOSITE TRACKING SCORE AND COMPETITIVE SHARE

For a weekly dashboard, combine visibility, recommendation rate, citation rate, and competitive share of voice into one composite:

$$\text{AIS} = V \times \text{Rec} \times \text{Cit} \times \text{SOV}$$

where $V = \text{AIV}(B)$ from Definition 1, Rec is the share of appearances that are recommendations, Cit is the share that are source-backed, and SOV_B is brand B 's share of all tracked brands' appearances:

$$\text{SOV}_B = \frac{\text{appearances}_B}{\sum_j \text{appearances}_j}$$

e.g. across 100 category questions, if the four tracked brands appear 78, 63, 52 and 19 times respectively, the fourth brand's share of voice is $19/(78 + 63 + 52 + 19) = 19/212 \approx 8.96\%$, a genuinely competitive metric, unlike "we rank #3 for one keyword."

6.5: TESTING WHICH FACTORS ACTUALLY PREDICT VISIBILITY

With a panel of (query, brand, engine, date) observations and covariates $X_1, \dots, X_K$ (referring domains, structured-data completeness, review volume, etc.), fit:

$$\Pr(Y = 1 \mid X) = \sigma\left(\beta_0 + \sum_{k=1}^K \beta_k X_k\right), \quad Y = 1 \iff \text{brand appeared}$$

This estimates *association*, not causation, a positive $\hat{\beta}_k$ says a factor predicts visibility, not that moving it will change your own outcome. The difference-in-differences design in §6.3, run on your own intervention, is what upgrades a correlational finding into causal evidence.

7. THE SIX LEVERS

What Theorem 1's factors mean in practice

R

Relevance

How closely your content matches the actual questions people ask, not the keywords you rank for. Answer-first structure and full intent-cluster coverage raise this.

A

Authority

Independent, credible sources treating you as a trustworthy reference in your category, press, analysts, high-authority reviews.

E

Evidence

The same claim about your brand, independently repeated across press, reviews, forums and docs, corroboration, not repetition on your own site.

T

Topical coverage

Breadth across the full query universe Q , informational, commercial, comparative and problem-oriented questions, not just your brand name.

S

Semantic clarity

One canonical name, category and description everywhere you appear. Every naming variant adds disambiguation risk the model must resolve before it will name you.

F

Freshness

Facts (pricing, specs, positioning) kept current. Stale, contradicted claims lower a system's confidence in citing you at all.

8. THE PLAYBOOK

Turning Theorem 1 into a 90-day plan

WEEKS 1–2: BASELINE

Build your query universe Q (30+ real buyer questions) and measure $\hat{p}_{\text{before}}$ across ChatGPT, Gemini and AI Overviews. Hold out a comparable control set of topics per §6.3.

WEEKS 3–5: S & F FIRST

Fix semantic clarity and freshness first, the neglected levers with the largest marginal $\beta_k \sigma'(z)$ per Corollary 1. Unify naming; ship Organization/Product schema; correct stale facts.

WEEKS 6–9: R & T

Raise relevance and topical coverage: rewrite key pages answer-first, one direct claim per heading, covering the full intent cluster in Q rather than the head keyword alone.

WEEKS 10–13: A & E

Build authority and evidence: earn independent corroboration of the same core claims across press, reviews and forums. Re-run $\hat{p}$ on both treatment and control, and compute $\widehat{\text{DiD}}$.

9. SCOPE & CAVEATS

What this model is, and isn't

- $\sigma(z_B)$ in Section 4 is a stylized model built to be qualitatively robust, not a disclosed or reverse-engineered ranking function of any specific AI search product, those are proprietary and change frequently.
- Theorem 1 and Corollary 1 are proved *from the model's stated assumptions* ($\beta_k > 0$), which is a directional claim, "these factors don't hurt, and returns diminish", not a numeric guarantee for any one brand.
- Section 5's proportionality is explicitly flagged as a heuristic, not a derived result, unlike Theorem 1.
- Section 6.5's regression establishes association; only the controlled experiment in §6.3, run on your own site, supports a causal claim about your own intervention.
- Be skeptical of anyone selling a specific "AI schema" as a guaranteed unlock, the model here argues for a portfolio of ordinary, verifiable signals (relevance, authority, evidence, coverage, clarity, freshness), not a trick.

Kashif Naveed · White Papers · AI Search Visibility